



# Container Orchestration using Docker & Kubernetes

By Mahendra Shinde

Microsoft  
Partner

Silver Learning  
Gold Cloud Platform



Educate

Advise

Implement

Manage

# Agenda

Module 1: Working with Kubernetes (Architecture)

Module 2: Working with Kubernetes (Deployment)

Module 3: Kubernetes on Azure





**SYNERGETICS**  
— GET IT RIGHT —

# Module 1

---

## Working with Kubernetes (Architecture)

Educate

Advise

Implement

Manage

# Kubernetes : Introduction

***“Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.”***

- One of the most popular cluster management tool.
- Having Largest developer community.
- Architecture designed by Google.
- Backed by 10 years of container deployment experience



# Why Kubernetes

- Manage your Application, ignore the machines.
- Manage Applications
  - Where to run
  - When to Get, when to run, and when to discard
- Application Image usage
  - Image lifecycle management
  - Run image as application (containers)

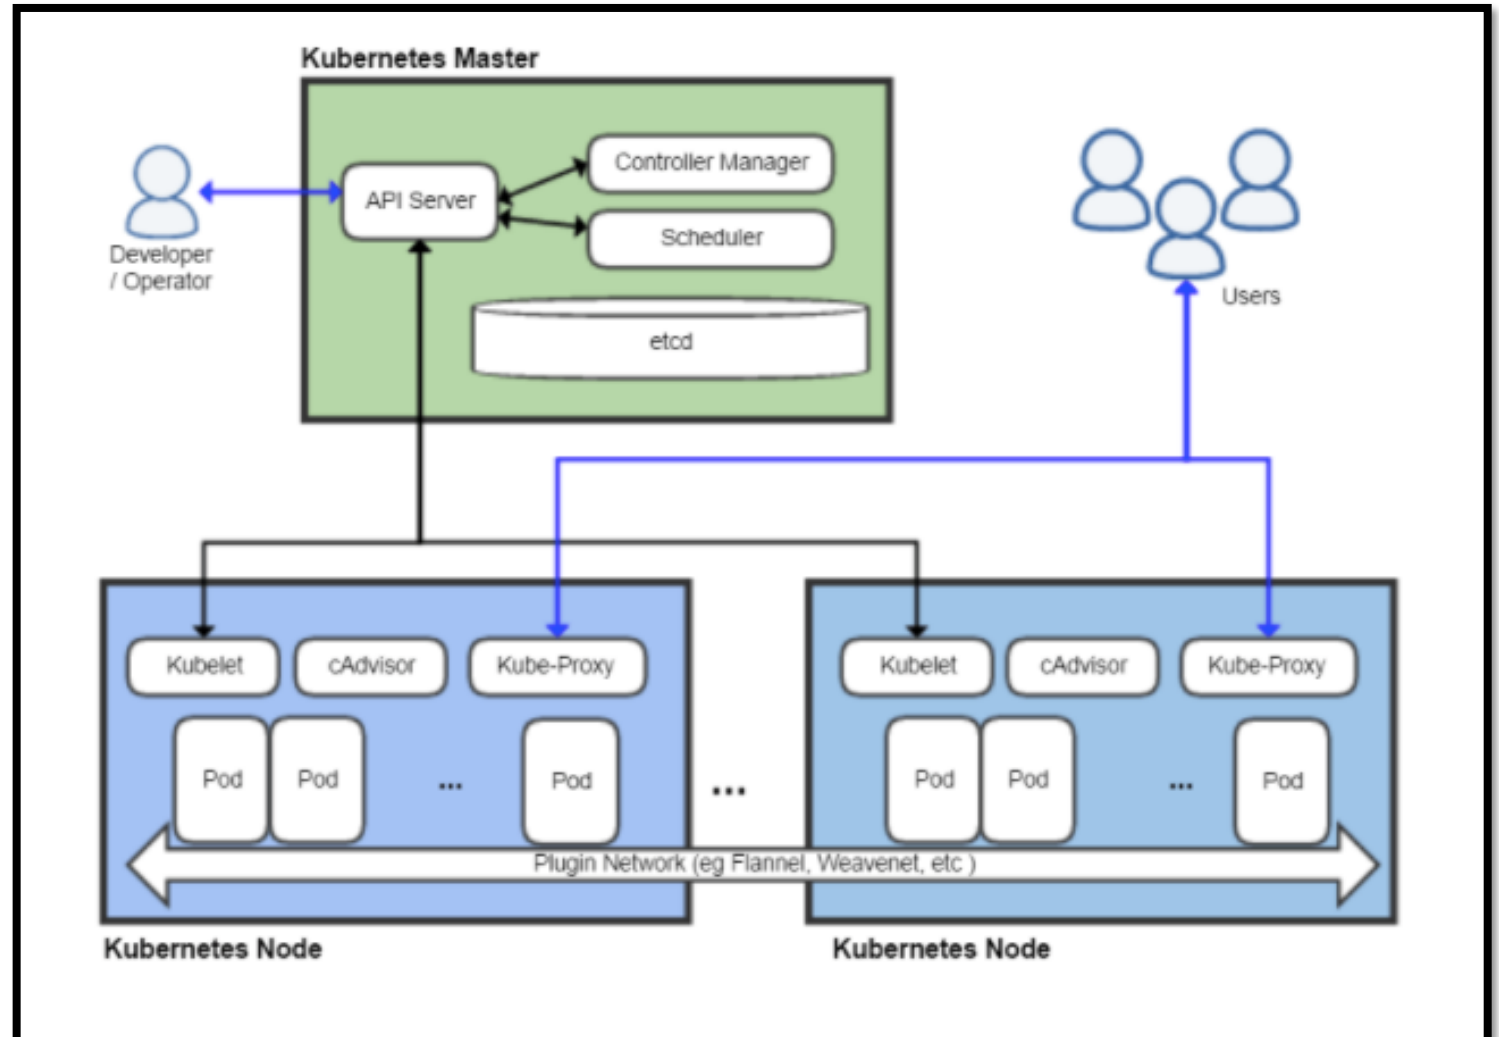


# What Kubernetes provide?

- High Level workload abstractions
  - If container resembles **atom**, kubernetes provide **molecules**
- Storage
- Network
- Monitoring
- Scaling
- Communication



# Kubernetes Architecture



# Kubernetes Components

## ■ Nodes

- A Physical or Virtual machines on kubernetes cluster.
- Can be either a master (manager) node or worker (deployment target) node.
- All the nodes must be in a private network
- Master nodes must have linux based operating system.
- Worker nodes can be either linux or windows based systems.
- Master node need few additional components than worker nodes





# Kubernetes : Master Components

- api-server
  - A Front end for kubernetes control plane.
  - Designed to scale horizontally.
- Etcd
  - Consistent & Highly available key-value store
  - Used as data store for all cluster data
- Scheduler
  - Put a pod on node!
  - Scheduling decision is based on
    - Affinity
    - Anti-affinity
    - Policy/Resource/Hardware/Software constraints
- Controller Manager
  - Responsible for following managers
    - Node controller
    - Replication controller
    - Endpoint Controller
    - Service Account and tokens controller



# Kubernetes : Node (Worker) Components

- Kubelet
  - An agent that runs on each node.
- Kube-proxy
  - Enables network abstraction by maintaining network rules on the host
  - Does connection forwarding
- Container runtime
  - Container runtime either Docker or RKT
  - Any other container runtime based on runc & oci
- Add-Ons
  - Cluster DNS
  - Web-UI



# Kubernetes Client tool (kubectl)


- CLI for managing kubernetes cluster
- Syntax :

`$ kubectl [command] [TYPE] [NAME] [flags]`

- Command: An operation like create, get, describe or delete
  - TYPE: A resource type like pod, service etc  
Case Sensitive, can be plural, like pods instead of pod
  - NAME: Name of resource to be operated
  - Flags: Optional flag, depends on resource type
- Example:
    - `$ kubectl get pods`
    - `$ kubectl describe deployment app1`
    - `$ kubectl get pod -f mypod.yml`



# Kubernetes Dashboard [Web-UI]



# Single Host deployment for developers

- Minikube :  
A simple standalone dev cluster for anyone who wants to get started with kubernetes.
- Uses Hypervisor like Oracle VirtualBox or HyperV.
- On Linux host, hypervisor could be set to NONE.



# Minikube architecture


- Entire cluster as single VM
- Host system as user workstation
- Use kubectl to access cluster

- Simple commands:

```
$ minikube start -vm-driver hyperv
--hyperv-virtual-switch="Swit
```

```
$ minikube delete
```

```
$ minikube dashboard
```



**A:** Minikube generates kubeconfig file

**B:** Minikube creates Minikube VM

**C:** Minikube sets up Kubernetes in Minikube VM

**D:** Kubectl uses kubeconfig to work with Kubernetes

# Multi Host deployments

- Ideal for production like environments
- Provides High availability [ by multiple master nodes ]
- Load balance workload across multiple worker nodes
- Secure & reliable cluster experience
- Minikube is ideal only for first time users to learn kubernetes
- Kubernetes provide following options
  - On cloud ( By vendors like Microsoft Azure, AWS, Google)
  - On Premise (Using kubeadm)



# kubeadm

- Kubeadm is tool to bootstrap kubernetes cluster (either on premise or on cloud)
- Uses Signed certificates for secure communication across cluster.
- Most common operations
  - Kubeadm init ( bootstrap a cluster from Master node)
  - Kubeadm join ( Join existing cluster from Worker node)
- System requirement
  - One or more machines running Debian or RedHat based linux
  - 2 GB or more RAM per machine
  - 2 or more CPUs per machine
  - Full network connectivity between all machines (Private/Public Network)







**SYNERGETICS**  
— GET IT RIGHT —

# Module 2

---

Working with kubernetes (Deployment)


Educate

Advise

Implement

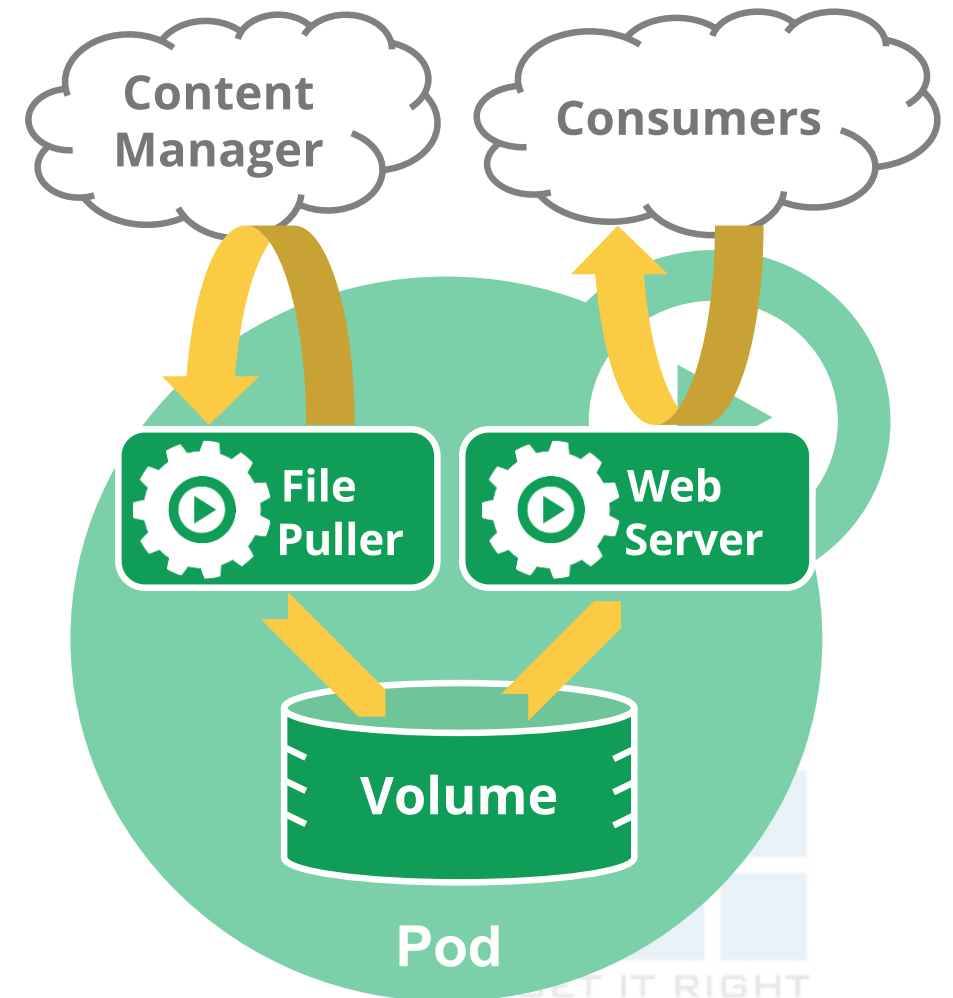
Manage

# Kubernetes Concepts



# Kubernetes Concepts : pods

- Smallest deployable unit for kubernetes.
- Ephemeral, disposable entities (ref cattle vs pets)
- May contain one or more container
- Deployed / destroyed / scaled as single unit
- All containers in a pod are sharing resources [Storage, Compute & Network]
- All container in a pod are tightly coupled.



# Kubernetes Concepts : pods

| Phase     | Description   |
|-----------|---|
| Pending   | Waiting for image downloading, not all containers ready yet.  |
| Running   | Pod deployed to a node, either all containers started or one is starting (images downloaded in pending phase)   |
| Succeeded | All Containers in the Pod have terminated in success, and will not be restarted.  |
| Failed    | All Containers in the Pod have terminated, and at least one Container has terminated in failure. That is, the Container either exited with non-zero status or was terminated by the system. |
| Unknown   | For some reason the state of the Pod could not be obtained, typically due to an error in communicating with the host of the Pod.  |



# Kubernetes Concepts : ReplicaSet

- A ReplicaSet ensures specified number of pod replicas
- A Replication controller makes sure that “Desired” replica count matches with “Actual” replica count for given ReplicaSet.



# Kubernetes Concepts : deployment

- A Desired State of Services, Pods, ReplicaSet etc.
- Can be written using YAML or JSON syntax
- Comparable to “docker-compose.yml” used in day 1
- Even deployments initiated by kubectl command are internally stored as YAML files.
- The sample deployment named “nginx-deployment”

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

# Kubernetes Concepts : service

- An abstraction which defines logical set of pods and policy to access them.
- Services establish a single endpoint for collection of replicated pods, distributing inbound traffic based on label selectors.
- In kubernetes modelling language they represent a “Load Balancer”.
- Pods and Services exists independently, have disjoint lifecycle.
- Supports both tcp & udp protocols (default: tcp)
- Kubernetes assigns VIP to service.
- Kubernetes DNS scheduler should assign a DNS label to each service.
  - Eg: service named “db” can be accessed by other services by name “db”



# Kubernetes Concepts : Secrets

- Intended to hold sensitive information like passwords, OAuth tokens or SSH keys
- Kubectl has few command to manage secrets

\$ kubectl get secrets

\$ kubectl create secret

A reference to secret can be added to deployment [yaml] file

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mypod
    image: redis
    volumeMounts:
    - name: foo
      mountPath: "/etc/foo"
      readOnly: true
  volumes:
  - name: foo
    secret:
      secretName: mysecret
```



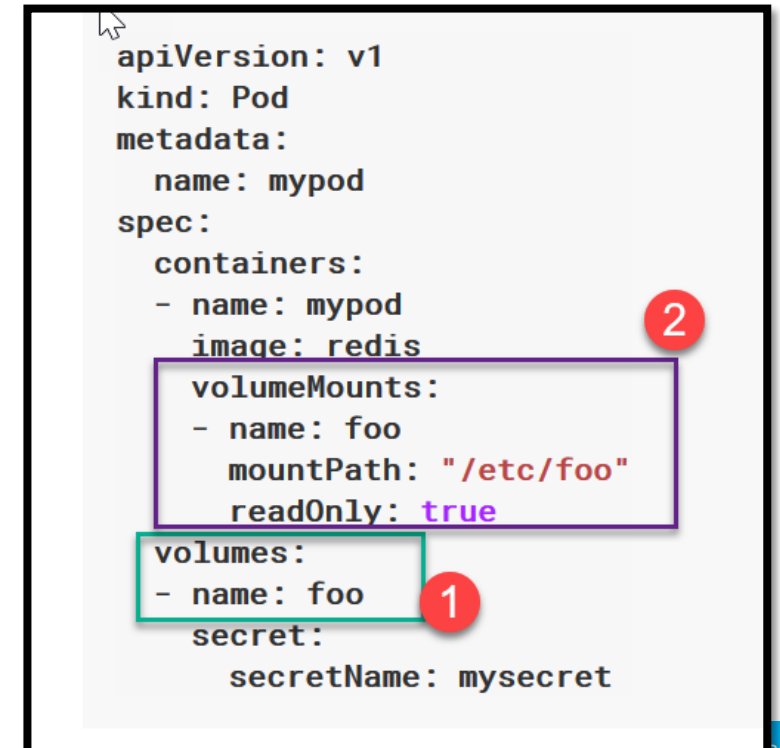
# Kubernetes Concepts : Volumes

- A Persistent data store for containers
- A container restart results in creating “clean container” (data in Writable layer lost!)
- Volumes can persist data between services or pods (containers sharing single pod)
- Deep down, Its just a directory made accessible to one or more containers / pods.
- Kubernetes supports following volumes
  - azureDisk
  - azureFile
  - Cephfs
  - Csi

```

apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mypod
    image: redis
    volumeMounts:
    - name: foo
      mountPath: "/etc/foo"
      readOnly: true
  volumes:
  - name: foo
    secret:
      secretName: mysecret

```



# Kubernetes Concepts: Networking

- Networking in kubernetes cluster is different than that of docker
  - Highly coupled Container to Container communication inside a pod. (use localhost)
  - Pod to Pod communication (Pods have IP address)
  - Pod to service communication (Services have VIP)
  - External to Service
- Implementations
  - ACI
  - Cilium
  - Contiv
  - Flannel
  - Weave Net



# Demos

- Demo : Deploying a sample application on kubernetes cluster



# Advanced Concepts

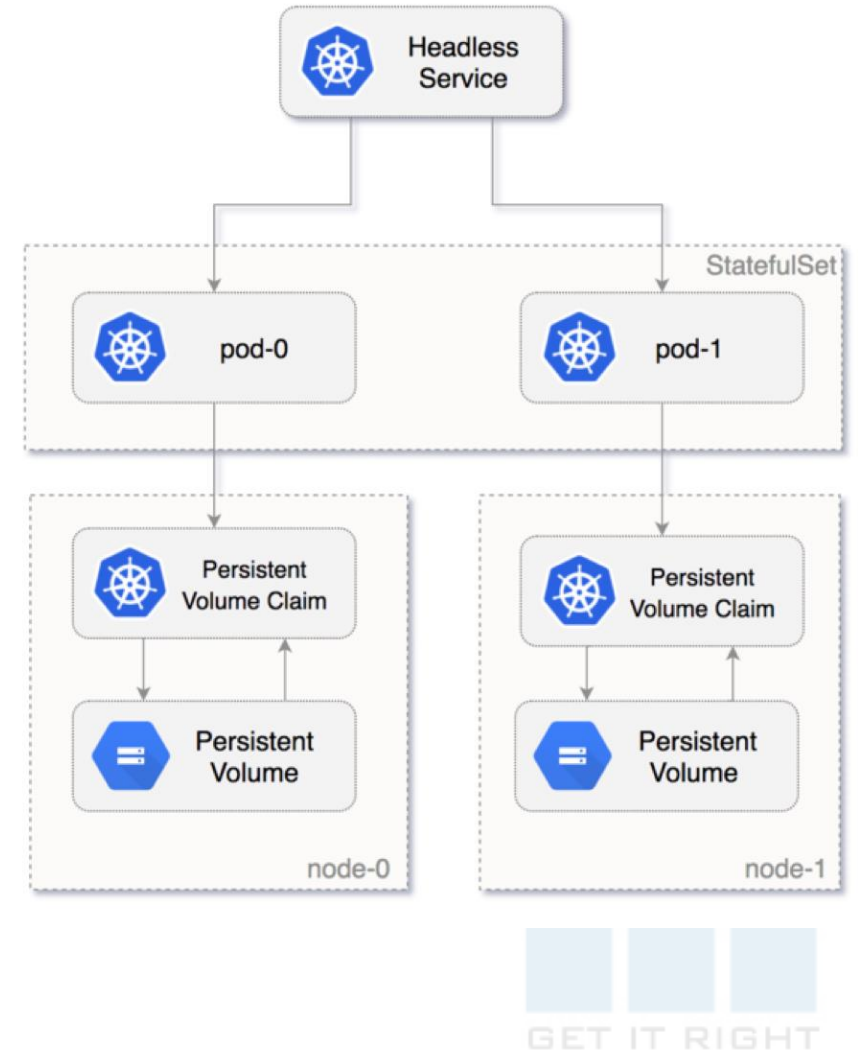
- Storage
  - Volumes
  - Persistent Volumes : Independent of pods, services or deployments
  - Storage classes (Used with Persistent Volumes)
- Compute
  - Pods can optionally define amount of resources like CPU or MEMORY
- Security
  - A Pod security policy (cluster level resource)
  - Can be managed by admission control plugin

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: azurefile
provisioner: kubernetes.io/azure-file
parameters:
  skuName: Standard_LRS
  location: eastus
  storageAccount: azure_storage_account_name
```

# Advanced Concepts

- Stateful Sets

- Targeted for Stateful applications
- Ordered Pod creation
- Stable network identity to each pod
- Uses PersistentVolumeClaims



# Advanced Concepts

- Scaling (Manual / Automatic)
  - Kubernetes uses “ReplicaSet” to ensure desired number of instances.
  - Manual scaling requires updating “replicas” property.
  - Auto (Horizontal) scaling can be created using “kubectl” command

```
$ kubectl autoscale rs app1rs --min=2 --max=5 --cpu-percent=80
```



# Advanced Concepts

## Rolling Update

- Update a service without an outage.
- Its works by
  - Creating a new replication controller with the updated configuration.
  - Increasing/decreasing the replica count on the new and old controllers until the correct number of replicas is reached.
  - Deleting the original replication controller.

- Example:

```
$ kubectl rolling-update NAME [NEW_NAME] --image=IMAGE:TAG
```

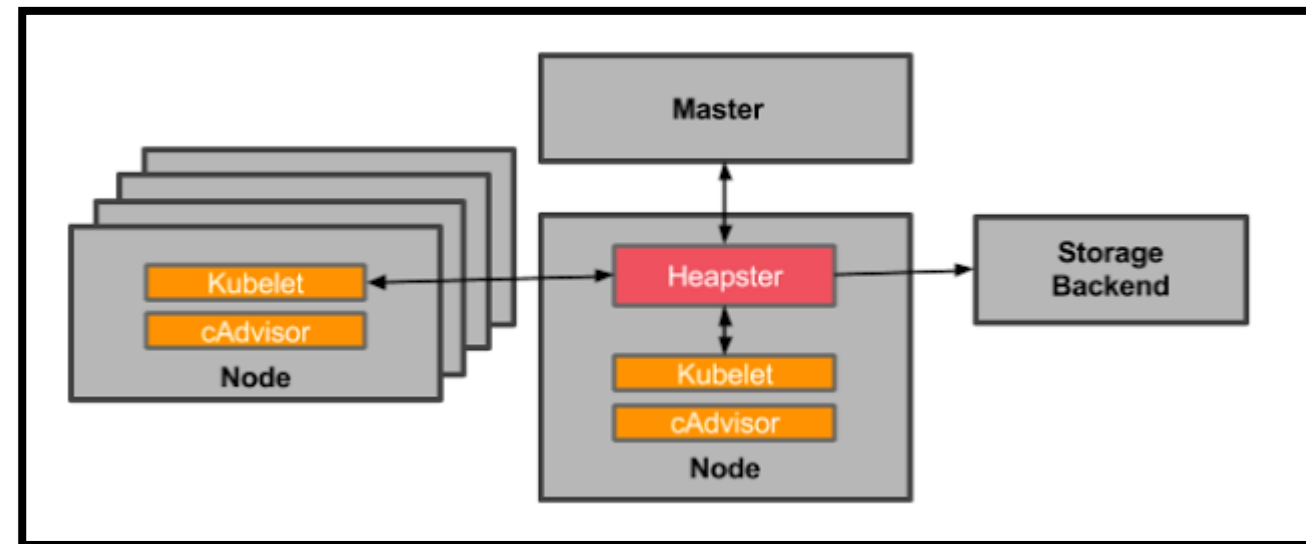


# Advanced Concept

## Monitoring (Infrastructure & Applications)

- Kubernetes dashboard provides basic monitoring
- Addon called “Heapster” can provide monitoring for entire cluster
- Uses cAdvisor

Container resource  
monitoring (native to docker)





# Advanced Concepts

## Disaster Recovery

- etcd is a backing store for entire cluster data.
- All data is stored as JSON objects
- Always have a backup plan for etcd
- Run etcd as cluster of odd numbers
- Restore etcd data into new cluster





**SYNERGETICS**  
— GET IT RIGHT —

# Module 08

---

## Kubernetes on Azure

Educate

Advise

Implement

Manage

# Azure Kubernetes Services (AKS)

- Azure has provided an alternative to ACS which uses “Kubernetes”.
- Easier to setup than ACS
- AKS is a managed service with a hosted Kubernetes control plane
- Has been certified as Kubernetes conformant.
- Is compliant with SOC, ISO, and PCI DSS.
- Differs from ACS in
  - Uses “Managed Disks” for all nodes
  - Currently support only one agent pool



# Hands On

- HOL 1: Deploying an AKS Cluster and accessing via kubectl in host system.
- HOL 2 : Deploying a sample application to AKS





**SYNERGETICS**  
— GET IT RIGHT —

# End Note

---

Educate

Advise

Implement

Manage



**SYNERGETICS**  
— GET IT RIGHT —

# Q/A

---

Thank You

Educate

Advise

Implement

Manage